

Lab 06. Deep Learning

Introduction to Computer Vision, Lab 06.

Today

- **Image classification**
- Recognition

Image classification

- Linear Model
 - A linear model performs a single linear transformation and can only represent linear relationships.

$$y = Wx + b$$

- In image classification, this becomes a major limitation because a **linear model cannot capture the complex, nonlinear patterns present in images**.

Image classification

- Multi-Layer Perceptron

- Multiple linear layers + nonlinear activation function.
- MLP works well on small, simple image datasets, e.g., MNIST digits where spatial structure is less critical.

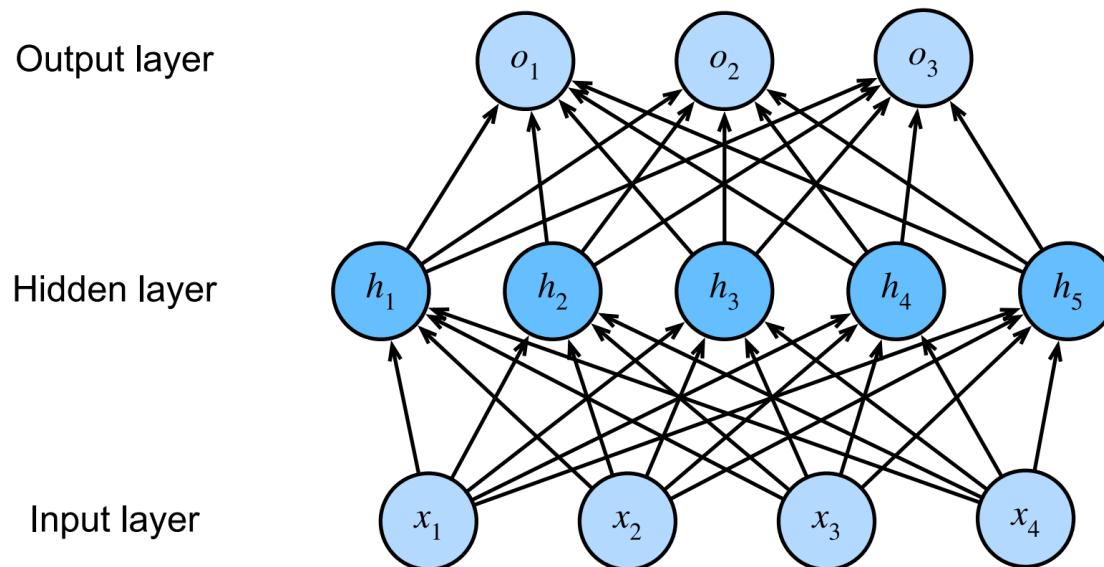


Image classification

- Convolutional Neural Network
 - CNNs take advantage of the **locality** of image data by using small receptive fields that capture local patterns.
 - CNNs achieve **translation invariance** through weight sharing, allowing them to detect the same features regardless of where they appear in the image.



Image classification

- Convolutional Neural Network
- See interactive CNN demo:
<https://poloclub.github.io/cnn-explainer/>

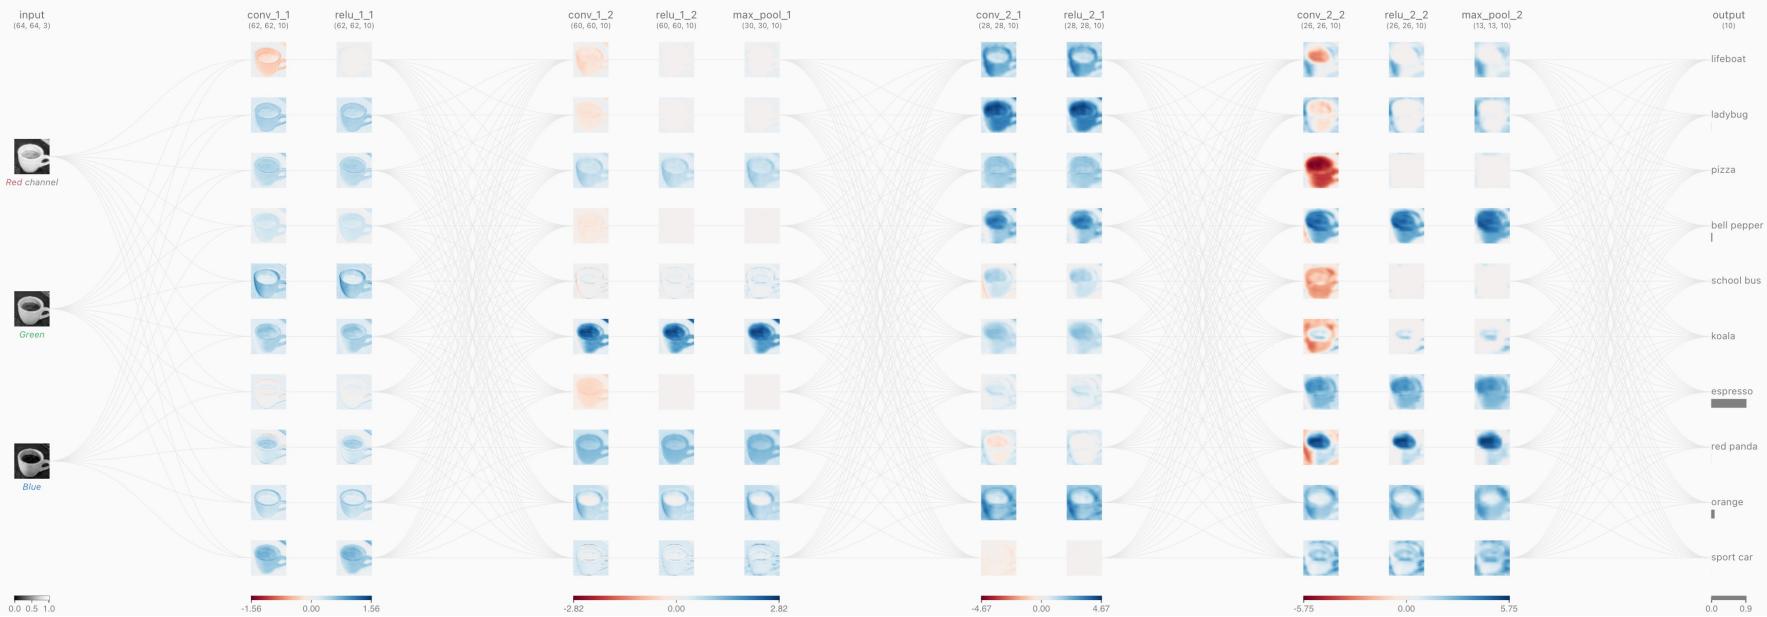


Image classification

- CIFAR10 dataset

- The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class.
- There are 50000 training images and 10000 test images.
- <https://www.cs.toronto.edu/~kriz/cifar.html>

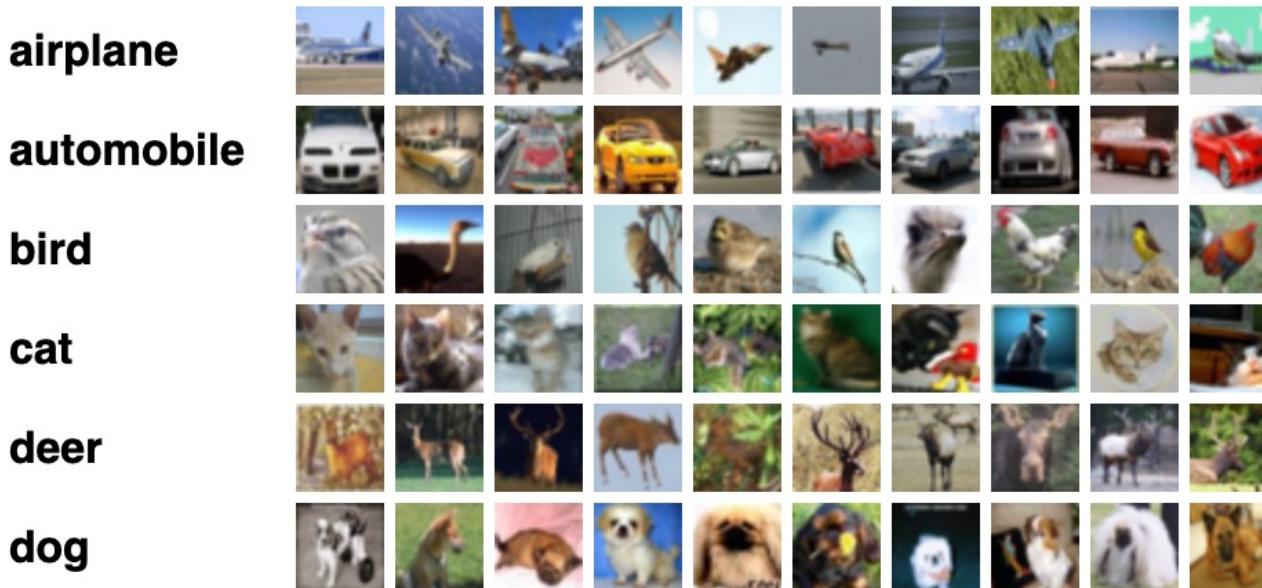


Image classification

- Please see classification.ipynb
- Train an MLP-based image classification model on CIFAR10 dataset.
- Build a CNN using PyTorch, train it and test it on CIFAR10 dataset.
- Add BN layers to the model, train it and test it on CIFAR10 dataset.

Pytorch Tutorial

- Please refer to slides from 2021 cs231n
Lecture 06: [2021 cs231n lec06.pdf](#)

Today

- Image classification
- **Recognition**

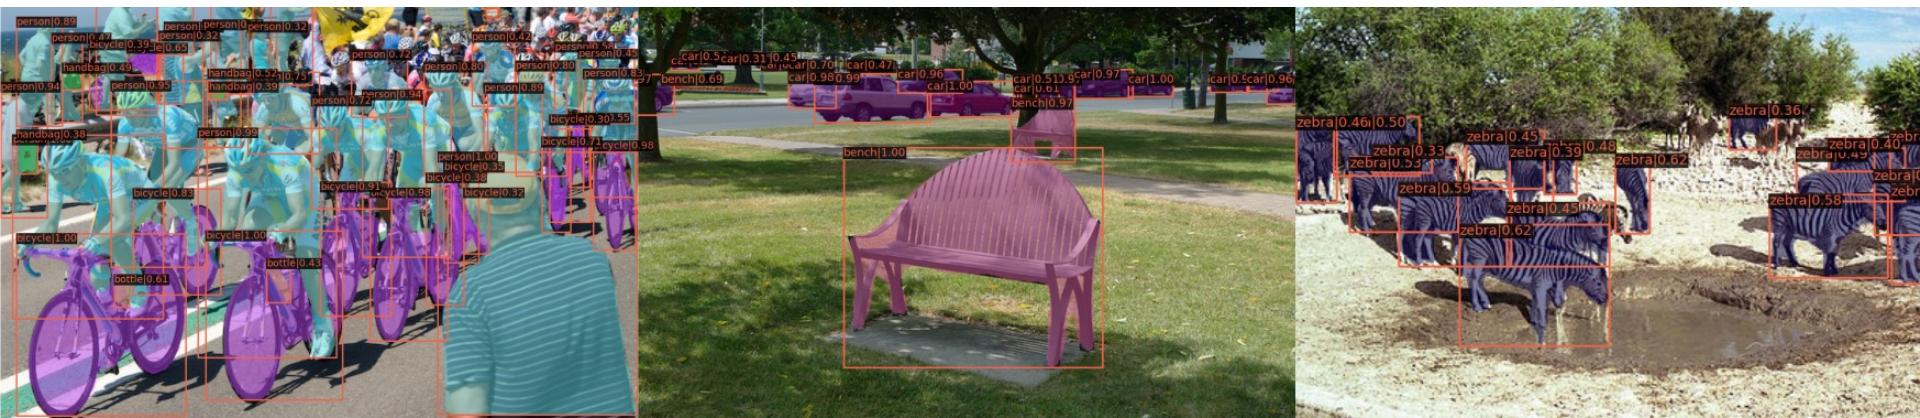
Recognition

- MMDetection



MMDetection is an open-source object detection and instance segmentation toolbox developed by OpenMMLab:

- **Modular design:** backbone, neck, head, loss, dataset, and pipelines can be freely combined.
- **Supports many classic and state-of-the-art models**, such as Faster R-CNN, YOLO series, RetinaNet, Mask R-CNN, and more.
- **Unified training/validation/inference workflow.**
- **Extensive model zoo** with many pretrained weights.



Recognition

- Popular 2D Detection Models
 - YOLOv3 / YOLOX
 - Fast RCNN / Faster RCNN
 - Mask RCNN / Cascade R-CNN

Recognition

- MMDetection Installation

- Latest:

https://mmdetection.readthedocs.io/en/latest/get_started.html

- v2.20.0:

<https://docs.qq.com/slide/DR0FnSGpxSkp4bGhv>

Tip: We provide two MMDet version examples of to accommodate different CUDA versions. You can choose **ANY MMDet version** which matches your local environment.

Recognition

- MMDetection Usage

- Download the pretrained model

```
mim download mmdet --config rtmde_tiny_8xb32-300e_coco --dest checkpoints/
```

- Run inference

```
from mmdet.apis import init_detector, inference_detector

config_file = 'configs_mmdet3.3/rtmde_tiny_8xb32_300e_coco.py'
checkpoint_file = 'checkpoints/rtmde_tiny_8xb32-300e_coco_20220902_112414-78e30dcc.pth'

model = init_detector(config_file, checkpoint_file, device='cpu')
img_path = 'pictures/det.jpg'
result = inference_detector(model, img_path)
```

Recognition

- Please see [recognition_mmdet3.3.ipynb](#)
- Running the MMDetection demo.
- Use your custom image.
- Use another model provided by this toolbox.

Submission

- Please only submit the two Jupyter notebooks.
- **Do not** submit any checkpoints or datasets.